

Where the ACL-2026 human-agent taxonomy meets this research, and where it stops

Author-generated. The same content is delivered as a standalone slide, VEGO_AI_ACL_Taxonomy_Slide_20260825.pptx.

What the taxonomy covers

01 Human Feedback

02 Communication Mode

03 Interaction Variant

04 Orchestration

4 established branches

The asymmetry is the argument about taxonomy coverage: four interaction branches versus eleven proposal-derived concepts.

What this proposal needs but the taxonomy cannot express

- 01 Reuse of a stored judgment in a later, different episode, and the reuse mode - inert, advisory, or
- 02 Claim-level validity scope: the prospective applicability envelope, including explicit negative scope
- 03 Diagnostic attribution: whether an intervention reveals a domain-specific quirk or a transferable capability gap
- 04 Temporal validity: expiry, supersession, revocation, and lapse when the interpreted guideline is revised
- 05 Claim-scoped authority and competence, separating case-level decisions from rubric-level changes
- 06 Version-exact provenance binding to the artifact state judged, with staleness detection
- 07 The elicitation trigger as a versioned, reason-coded policy object
- 08 Attention-budget accounting: a bounded budget per run and an allocation rule across claims
- 09 Preserved dissent: two conflicting judgments both retained, reuse blocked pending adjudication
- 10 Reuse-leakage control: provenance disjointness between the judgment store and the evaluation cases
- 11 Judgment target layer: verdict, the agent's stated reasoning, evidence selection, or the guideline

This states what the taxonomy cannot express - not that those concepts are necessary or that the proposed mechanisms work.