

VEGO-AI PhD — Supervisor Package, 12 August 2026

Prepared by: Ali Hamed · **For:** Prof. Iris Reinhartz-Berger and Prof. Arnon Sturm **Status:** working materials for discussion — **not approved, not final.** The research-question wording in this package is Ali's reconstruction of the live edits from the 5 August working call, and confirming or correcting it is the first item on the agenda.

What is in this folder

#	File	What it is
01	Executive Brief	One page, English and Hebrew — the whole picture in about a minute
02	Chapter 3 — Gap and Research Questions	The written chapter requested on 5 August: the research gap, the umbrella question, and SQ1–SQ3 with the reasoning behind each
03	Literature Review — Per Research Question	The per-question literature spreadsheet requested on 5 August: every source tagged RQ1/RQ2/RQ3/general, plus an honest coverage-gap view
04	Progress Presentation	The slides for the 12 August meeting (PowerPoint, with a PDF copy)
05	Requirements and Progress Tracking	Every requirement, action, and open question from both supervision calls, with current status — and the ten decisions requested at this meeting

Where things actually stand

Written and ready for your review: the Gap & Research Questions chapter is drafted in full and incorporates each correction from the 5 August call. The literature spreadsheet is built and tagged per research question. The four-agent architecture and its supporting mechanisms are implemented and tested.

Deliberately not done yet: the literature searches themselves have not been run, and the per-question artifact design has not been written — both are on hold per the instruction on 5 August to think about sections 2 and 4 without executing them yet.

Not yet established, and not claimed anywhere in this package: no accuracy, generalization, or effort-reduction result exists. The independent expert-labelling gate stands at 0 labels supplied out of 24 generalization-safe candidate rows, against a required minimum of 20 — so those measurements are not yet computable. Medical-track readiness stands at 0 of 6 entry gates, and no clinical claim of any kind is made.

What is being asked of you at this meeting

Six decisions are needed: the umbrella research-question wording; the SQ1–SQ3 wording; the three-study mapping and method direction; the Plan A / Plan B boundary; confirmation that every research question stays answerable under Plan B; and the evidence-boundary wording. Four further items can be settled if time allows, including the fallback control date and the initial literature scope. All ten are listed in file 05, sheet "Decisions Requested."

A note on sourcing

Material attributed to the 5 August discussion is drawn from a machine-generated transcript in which speakers are inferred rather than confirmed. Those attributions are the record's inference and remain subject to your correction. Nothing in this package quotes either of you verbatim.