

Chapter 6 - Work Plan

Working draft. Semester-aligned blocks over a 3-year horizon per E14. Dates are internal working targets, not confirmed university deadlines.

Status: complete working draft. Every research-deliverable date in this chapter is an internal working target unless explicitly marked otherwise. The university process (action A-14) is open: no submission, candidacy, or review date has been confirmed in writing by an authoritative university source, and the September/October 2026 targets carried by earlier planning documents remain unverified. This chapter is written so that a confirmed university calendar can be dropped into it without changing the research sequence.

6.1 Planning principles

6.1.1 Blocks, not months

The plan is expressed in **approximately three-month blocks** across a three-year horizon. It is deliberately not expressed month by month, and this is a methodological commitment rather than a drafting convenience.

Three of the four things that gate this research have latency that the candidate does not control:

- **independent expert returns** — the evidence base for every quantitative claim depends on named human reviewers completing blinded labelling and adjudication on their own schedule;
- **institutional decisions** — supervisor sign-off on question wording and study protocols, university process confirmation, and (under Plan A) ethics, data-custodian, and environment approvals;
- **external review** — paper submission-to-decision cycles.

A month-level schedule would assign precise dates to activities whose start conditions are held by other people. That produces a plan that is wrong within weeks and, worse, a plan that invites the reader to treat unverified dates as commitments. A block-level schedule states what must be *true* by the end of a roughly three-month window, names the gate that must clear first, and leaves the internal ordering to the weekly supervision cadence (§6.5.1), which is the correct instrument for week-level control.

The trade is explicit: the plan gives up apparent precision in exchange for being honest about which of its dates are load-bearing. Only the gates in §6.3 are load-bearing.

6.1.2 Why semester alignment

The blocks align to the academic calendar rather than to calendar quarters, for four reasons that bear directly on feasibility:

1. **The evidence source is course-generated.** The motivating corpus is student-produced domain models from taught courses. Any extension of the corpus, any second course setting under Plan B, and any new cohort of models becomes available on a teaching-semester rhythm, not a calendar-quarter rhythm.
2. **Expert reviewers are academics.** Independent labelling, adjudication, and second-panel recruitment compete with teaching load. Blocks that straddle a teaching semester and its examination period will systematically under-deliver on human-dependent tasks; blocks that place labelling campaigns in inter-semester windows will not.
3. **Supervision intensity varies by term.** Both supervisors carry teaching and examination duties. The weekly cadence continues year-round, but the depth of review realistically available differs between a teaching block and an inter-semester block, and design-freeze decisions are scheduled into the latter.
4. **Institutional milestones are semester-indexed.** Progress reporting, candidacy scheduling, and committee availability follow the academic year. Aligning the plan to it means that when A-14 is answered, the answer maps onto an existing block boundary instead of forcing a re-plan.

The precise term dates at the University of Haifa and Ben-Gurion University have **not** been verified for the years covered here. The block grid is therefore stated at quarter resolution — a two- or three-week shift in term boundaries does not move any block.

6.1.3 Date status vocabulary

Every date in this chapter carries exactly one of four statuses. This vocabulary is used consistently and should be preserved in any revision.

Status	Meaning	Example
Verified	Confirmed by an external, authoritative source, with the evidence recorded	Weekly supervision series, Wednesday 09:00–10:00 Asia/Jerusalem, accepted by Ali, Iris, and Arnon through the 2026-10-07 occurrence (DC-008)
Internal working target	A date this project set for itself; no external party is committed to it; it may move without any external consequence	Proposal submission in the autumn of 2026
Contingent	Cannot be dated at all until a named gate clears; the block position shown is provisional	Any Plan A medical activity (0 of 6 gates passed)
Structural	Fixed by the academic calendar, independent of this project	Block boundaries

The single rule that governs the whole chapter: **until action A-14 produces a written answer from an authoritative university source, every submission, candidacy, and review date in this plan is an internal working target.** The University Process Verification Checklist currently records all eleven required fields — official deadline, reviewer count and eligibility, nomination process, committee composition, forms and submission channel, length and formatting, candidacy presentation requirements, required lead times, whether a preliminary-results section is mandatory, the amendment process, and the authoritative policy source — as *Pending*, with the owner unassigned. If a written answer establishes an earlier binding deadline, the grid below rebaselines within one working day, preserving the research sequence and the gate structure but compressing or re-ordering the blocks.

6.1.4 The block grid

Block	Window	Season label	Character
Block 0	Aug–Sep 2026	Bridge (part block)	Pre-grid; ~7 weeks, not a full block
Y1 B1	Oct–Dec 2026	Autumn	Teaching semester A
Y1 B2	Jan–Mar 2027	Winter	Semester A close, examinations, inter-semester
Y1 B3	Apr–Jun 2027	Spring	Teaching semester B
Y1 B4	Jul–Sep 2027	Summer	Inter-year
Y2 B1	Oct–Dec 2027	Autumn	Teaching semester A
Y2 B2	Jan–Mar 2028	Winter	Inter-semester
Y2 B3	Apr–Jun 2028	Spring	Teaching semester B
Y2 B4	Jul–Sep 2028	Summer	Inter-year
Y3 B1	Oct–Dec 2028	Autumn	Teaching semester A
Y3 B2	Jan–Mar 2029	Winter	Inter-semester
Y3 B3	Apr–Jun 2029	Spring	Teaching semester B
Y3 B4	Jul–Sep 2029	Summer	Closing block and schedule buffer

Twelve full blocks span October 2026 to September 2029 — the three-year horizon confirmed in the 2026-08-05 supervision discussion (attributed to Arnon in the machine-derived record, which corrected an earlier two-year assumption). Block 0 is shown separately and is explicitly *not* a full block: it covers the current position of the work at the time of writing and is roughly seven weeks long.

Two notes on the numbering. First, "Year 1" here denotes the first full academic year covered by the plan (2026/27); whether the university counts doctoral years from registration, from proposal approval, or from candidacy is part of A-14 and is unverified. Second, resolution decreases deliberately across the grid. Blocks 0 through Y1 B4 are specified against work that is concretely pending today and can be described precisely. Year 3 blocks are specified at the level of objective and exit criterion only. Writing detailed deliverables for the spring of 2029 would be fabrication, and the plan says so rather than performing false confidence.

6.2 The block-by-block plan

6.2.1 Overview

Block	Primary focus	Advances
Block 0	Decision closure, proposal assembly, evidence mobilisation	All

Y1 B1	Literature review executed and synthesised; Study 1 design frozen	SQ1 (SQ2, SQ3 supported)
Y1 B2	Study 1 artifact built and validated; EXP-005 closed to its gate	SQ1, SQ3 evidence base
Y1 B3	Study 1 empirical evaluation; Paper 1	SQ1
Y1 B4	Study 2 design frozen; annual review; medical gate re-review	SQ2
Y2 B1	Governed judgment lifecycle artifact built	SQ2
Y2 B2	Study 2 empirical evaluation with independent experts	SQ2
Y2 B3	Paper 2; Study 3 preregistered	SQ2 → SQ3
Y2 B4	Transfer arm instantiated in the second setting	SQ3
Y3 B1	Study 3 evaluation and transfer classification	SQ3
Y3 B2	Paper 3; thesis integration begins	SQ3, U-RQ
Y3 B3	Thesis assembly and internal review	U-RQ
Y3 B4	Revision, submission, defence preparation; buffer	—

6.2.2 Block 0 — Bridge (August–September 2026, part block)

Objective. Convert the current provisional state into a submittable proposal, and close the two decisions on which every downstream block depends.

Deliverables.

- A logged decision on D-RQ-01 (umbrella research question wording) and D-RQ-02 (SQ1–SQ3 wording), preceded by verification of the live wording against the candidate's own saved working draft from the 2026-08-05 call. Until this is logged, the wording carried in Chapter 3 remains provisional.
- Chapter 6 (this chapter) and Chapter 5 (Preliminary Results, software/modelling only) drafted; Chapter 3 already complete; Chapter 1 (Introduction) deferred to last, per the confirmed chapter order.
- Chapters 2 (Literature Survey) and 4 (Methodology) released from their current deferral and drafted. *These two chapters are presently deferred by supervisor instruction (A08-04: think about, do not start). The plan assumes release at or shortly after the 2026-08-12 review. If release is deferred further, the submission target moves — and because that target is an internal working target, nothing external breaks.*
- The **2026-08-26 internal checkpoint** executed and its outcome logged: for each of the six medical entry gates G1–G6, either a documented owner, evidence path, and feasible completion date, or a recorded failure. This date is an internal project-control date (DC-005), **not** a university deadline and not attributable to any supervision meeting.
- The A-14 university-process inquiry assigned an owner, authorised, sent, and — if answered — recorded with the authoritative source, contact, and date verified.
- Reviewer recruitment for EXP-005 initiated under the approved labelling protocol: two independent reviewers plus an adjudicator, with the supervisor label-approval pack signed off before any blind sheet is released.

Advances. All three sub-questions, indirectly: this block removes the blockers rather than producing research evidence.

Evidence produced. Logged decision records with old wording, new wording, rationale, approver, and effective date; a logged checkpoint outcome with per-gate status; a written university-process response or an explicit record that none was received; a signed labelling approval pack.

Exit criteria.

5. D-RQ-01 and D-RQ-02 are logged as decided or explicitly corrected — not merely discussed.
6. The committed execution path (Plan A or Plan B) is recorded following the 2026-08-26 checkpoint.
7. A proposal draft exists containing all six chapters, with every claim carrying an evidence state and no claim exceeding the available evidence.
8. The university process is either verified in writing, or recorded as unverified with the schedule explicitly flagged provisional throughout the proposal.

Not in this block. No accuracy, effort, or generalization claim; no patient-row work of any kind; no assertion that the literature review is complete.

6.2.3 Year 1, Block 1 — Autumn (October–December 2026)

Objective. Turn the frozen literature protocol into an executed, screened, synthesised evidence base, and freeze the Study 1 design.

Deliverables.

- **Execution of QL-01 through QL-05** across their planned databases (ACM Digital Library, IEEE Xplore, Scopus, Web of Science; PubMed for QL-05), with the full execution record required by the protocol for every database row: exact executed query copied from the platform, date/time and timezone, filters and field restrictions, total hits, export filename and SHA-256, platform transformations, deduplication batch, and screening counts with controlled exclusion reasons. Execution follows the RQ sign-off because the queries encode the question concepts.
- Title/abstract screening, full-text screening, deduplication, and identity verification, with one paper per row in the literature workbook and the authors' conclusions kept separate from the candidate's synthesis.
- **Chapter 2 completed from executed searches**, replacing the proposal-stage survey that drew on the existing 139-source workbook. Every literature-gap statement in Chapter 3 §3.1 is currently a *candidate claim*; this block either establishes each one or corrects it, and Chapter 3 is revised accordingly.
- **Study 1 design frozen and preregistered**: the event and uncertainty contract, intervention eligibility and priority criteria, the dosage and attention-budget model, routing and escalation/timeout rules, the burden model, and the requirements-to-artifact traceability matrix.
- **EXP-005 mobilised**: calibration returns obtained from both reviewers, labelling instructions frozen by a recorded human decision, and the blinded evaluation set released only after that gate — with reviewer returns preserved immutably.

Advances. SQ1 primarily (Study 1 design); SQ2 and SQ3 through the literature synthesis and the evidence pipeline.

Evidence produced. A complete search execution register with hit counts and export hashes; screening counts and exclusion reasons; a revised gap chapter in which each claim is marked established, corrected, or refuted; a frozen Study 1 protocol; immutable reviewer-1 returns.

Exit criteria.

9. Every planned database row for QL-01–QL-05 has either an execution record or an approved, explained not-applicable disposition.
10. Chapter 2 is written from screened results, and any novelty claim in the proposal is now supported by an executed review rather than asserted.
11. The Study 1 protocol is frozen and accepted or corrected by the supervisors before any evaluation runs.
12. Reviewer-1 returns exist and are preserved; if they do not, the block exits with a recorded reason and the labelling campaign carries into Y1 B2 as the block's first commitment.

6.2.4 Year 1, Block 2 — Winter (January–March 2027)

Objective. Build and validate the Study 1 artifact, and close EXP-005 to its minimum evidence threshold.

Deliverables.

- The selective-intervention architecture constructed in the software/modelling instantiation: listener and event catalogue, eligibility and priority criteria, triage and dosage policy, human-routing contract, timeout and escalation rules, and the burden budget.
- Requirements-to-artifact traceability: every SQ1 requirement mapped to an inspectable artifact and a validation.
- Bounded scenario validation across controlled event, uncertainty, priority, dosage, routing, escalation, and burden scenarios. Synthetic fixtures are mechanism tests only and are labelled as such.
- **Reviewer-2 returns and adjudication**, producing a frozen gold set. The minimum evidence threshold is **at least 20 adjudicated generalization-safe labels**; the labelling protocol's preferred target is 30–50, which exceeds the 24 generalization-safe candidates currently available and therefore requires either a second labelling round or an extension of the candidate set — a decision to be taken with the supervisors in this block.
- The first quantitative analysis that is *eligible* under the evidence gate, on the development split only, with the holdout sealed.

Advances. SQ1 (artifact and validation); SQ3 (the evidence base for later transfer testing).

Evidence produced. An inspectable architecture with a complete traceability matrix; scenario validation results; an adjudicated gold set with inter-rater agreement where enough reviewer-2 labels exist; a preregistration of any candidate change made before the holdout is opened.

Exit criteria.

13. Every SQ1 intervention requirement maps to an inspectable artifact and a validation; burden and uncertainty controls are explicit.
14. Either at least 20 adjudicated generalization-safe labels exist and the gold set is frozen, **or** the exact block is recorded with its cause and the quantitative claims remain formally unavailable. Both outcomes satisfy this block; only silence about the state would fail it.
15. No effectiveness, accuracy, or workload claim has been made on the basis of architecture, tests, or fixtures alone.

6.2.5 Year 1, Block 3 — Spring (April–June 2027)

Objective. Run and report the Study 1 empirical evaluation; submit the first paper.

Deliverables.

- Controlled evaluation of intervention policies under an explicit attention budget, measuring coverage of important uncertainty against expert load — the two-sided criterion built into SQ1, under which a policy that achieves coverage by flooding the expert fails by definition.
- Error and validity analysis, including the threats register for Study 1.
- **Paper 1** submitted: a selective-intervention / agentic human-in-the-loop architecture paper to a human–AI interaction, agentic-systems, or AI/software-engineering venue.
- A results package in which every finding carries an evidence state and negative or null findings are reported rather than suppressed.

Advances. SQ1.

Evidence produced. Measured coverage-versus-burden results where the label evidence permits; the validity analysis; a submitted manuscript with its submission record.

Exit criteria.

16. The Study 1 exit criterion in the three-study contract is met, or its shortfall is stated precisely.
17. Paper 1 is submitted (submission, not acceptance, is the milestone — acceptance is outside the candidate's control and is therefore never used as a gate in this plan).
18. No result is reported beyond what the adjudicated label evidence supports.

6.2.6 Year 1, Block 4 — Summer (July–September 2027)

Objective. Freeze the Study 2 design; complete the first annual review; re-review the medical gates.

Deliverables.

- **Study 2 design frozen and preregistered:** the structured judgment schema (including the system's own core reasoning, not only the expert's verdict); source-grounded verification and challenge protocol; conflict, adjudication, and authority rules; the provenance model; scope controls, expiry, and revocation; explainable retrieval; and the bounded integration and reuse policy.
- A first-year progress report to the supervisors and, if the process requires one, to the department. *The existence, form, timing, and assessment of any departmental annual report are unverified and part of A-14.*
- **Annual re-review of medical entry gates G1–G6,** with a logged verdict. As of the time of writing the standing verdict is 0 of 6, and a further complication is on record: a third clinical dataset (EHRSHOT, OMOP-CDM patient-level tables) was observed in a supplied Drive folder on 2026-08-11 and is named in none of the gates, which were written against MIMIC and Clalit. Gates 1 and 3 both require a *named* data source, so this must be resolved before either gate can even be assessed. Nothing in that folder has been opened, copied, or processed.
- Carry-forward of any EXP-005 shortfall as an explicit, owned item rather than an assumption.

Advances. SQ2.

Evidence produced. A frozen Study 2 protocol with supervisor acceptance; a logged year-1 report; a per-gate medical verdict with named owners and evidence paths or recorded absences.

Exit criteria.

19. The Study 2 protocol is frozen before construction begins.
20. The G1–G6 verdict is logged, whatever it says.
21. Year-1 progress is accepted by the supervisors, with any correction recorded as a decision rather than silently absorbed.

6.2.7 Year 2, Block 1 — Autumn (October–December 2027)

Objective. Build the governed judgment lifecycle artifact.

Deliverables. The implemented judgment schema and its contract validation suite; source-grounded verification and challenge; conflict detection, reconciliation, and authority transitions; the provenance chain; explainable retrieval with complete retrieval explanations; scope, expiry, and revocation enforcement; and bounded reuse that remains advisory and gated.

Advances. SQ2.

Evidence produced. Schema and contract validity results; provenance completeness measures; validation and reconciliation coverage; unsafe-reuse rejection tests; scope, expiry, and revocation enforcement evidence.

Exit criteria. Every SQ2 lifecycle stage, authority rule, and unsafe-reuse control maps to an inspectable artifact and a validation. Transparent reuse remains advisory and gated; no claim is made that governed reuse improves outcomes on the strength of schemas, memory, tests, or fixtures.

6.2.8 Year 2, Block 2 — Winter (January–March 2028)

Objective. Evaluate Study 2 with independent experts.

Deliverables. A second, larger independent labelling and judgment-capture round with a recruited expert panel; a conflict and reconciliation study using controlled conflict fixtures alongside real judgments; unsafe-reuse rejection testing against the scope controls; leakage controls between the context that produced a judgment and the context in which it is evaluated; and a bounded-convergence analysis.

Advances. SQ2, with the resulting judgment corpus feeding SQ3.

Evidence produced. Independent expert returns preserved immutably; measured conflict and reconciliation behaviour; recorded rejections of out-of-scope reuse attempts.

Exit criteria. Study 2 measures are computed, **or** a recorded and reasoned statement explains precisely why each is not computable. Expert authority is demonstrably preserved: for every reused judgment, the originating expert can see, contest, and revoke it.

6.2.9 Year 2, Block 3 — Spring (April–June 2028)

Objective. Report Study 2; preregister Study 3.

Deliverables.

- **Paper 2** submitted: a governed knowledge-reuse and judgment-lifecycle paper to a human-centred AI, knowledge-engineering, or AI-assisted model-engineering venue.
- **Study 3 preregistered and frozen before any transfer data is opened**, including: the classification criteria separating *domain-specific* uncertainty from *broadly transferable* uncertainty — the analytic core of SQ3 following the correction recorded in the 2026-08-05 discussion (attributed to Arnon in the machine-derived record), under which failure to identify actors and use cases is a **general capability gap** whose judgments are transfer candidates, while judgments attached to domain-bound uncertainty must not travel; the leakage controls; the paired-comparison design; the workload measures; and pre-registered success and stop criteria.
- Selection and authorisation of the second setting: under Plan B, an authorised second software/modelling context (a different diagram family or modelling task, a second annotated corpus, an external course or institution, a different reviewer panel or expertise profile, or a longitudinal reuse setting); under Plan A, the approved clinical context, contingent on all six gates.

Advances. SQ2 completion, SQ3 initiation.

Evidence produced. A submitted manuscript; a frozen, timestamped Study 3 preregistration; a written authorisation for the second setting.

Exit criteria. The Study 3 protocol is preregistered *before* the second setting is opened. The transfer classification criteria are stated in a form that can be applied by someone other than the candidate and can be shown to be wrong.

6.2.10 Year 2, Block 4 — Summer (July–September 2028)

Objective. Instantiate and run the transfer arm; complete the second annual review.

Deliverables. The second setting instantiated with construct comparability preserved; reviewers recruited for that setting; independent labels collected under the same blinding and adjudication discipline; the sealed-holdout discipline maintained and the original baseline not overwritten; a second-year progress report; and a further logged G1–G6 re-review.

Advances. SQ3.

Evidence produced. Independent labels in the second setting; a provenance record for construct comparability; the year-2 report.

Exit criteria. The transfer arm is running with the holdout sealed and leakage controls demonstrably in place; the year-2 report is accepted; the medical verdict is logged.

6.2.11 Year 3, Block 1 — Autumn (October–December 2028)

From this point the plan is stated at objective and exit-criterion resolution. Detailed deliverables two years out would be invented.

Objective. Complete the Study 3 evaluation and the transfer classification.

Scope. Paired comparison across settings; the common-core versus adaptation map; classification of each captured judgment as domain-specific or broadly transferable against the preregistered criteria; adaptation effort; the ratio of invariant to domain-specific components; and the reasons transfer was blocked where it was blocked.

Advances. SQ3.

Exit criteria. The transfer taxonomy is validated against independent labels rather than asserted; null and negative transfer results are retained and reported; boundary conditions are fed back to the SQ1 and SQ2 policies as the three-question structure requires.

6.2.12 Year 3, Block 2 — Winter (January–March 2029)

Objective. Submit the third paper; begin thesis integration.

Scope. **Paper 3B** — a cross-domain or cross-institution software/modelling evaluation-and-transfer paper — or **Paper 3A**, a clinical-informatics evaluation-and-transfer paper, *only* if Plan A gates and evidence exist. Consolidation of the result and validity package across all three studies. First integrated thesis chapter drafts.

Exit criteria. Paper 3 submitted; every claim intended for the thesis is mapped to a specific evidence item and an evidence state.

6.2.13 Year 3, Block 3 — Spring (April–June 2029)

Objective. Assemble the thesis and complete internal review.

Scope. Full thesis draft; supervisor review cycles; a traceability audit in which every claim resolves to evidence and every cited artifact path resolves; a reproducibility package covering code, environment, parameters, runs, and artifact provenance.

Exit criteria. A complete draft accepted by both supervisors as ready for external examination, subject to whatever process A-14 establishes.

6.2.14 Year 3, Block 4 — Summer (July–September 2029)

Objective. Revise, submit, prepare the defence — and absorb accumulated slippage.

Scope. Revision against supervisor and reviewer comments; submission under the then-verified university process; defence preparation.

Exit criteria. Thesis submitted. This block is deliberately the least loaded in the grid: it is the plan's only structural buffer, and the risk analysis in §6.6 assumes it will be partly consumed.

6.2.15 What is deliberately absent from this plan

No block's exit criterion requires a positive result. No block requires accuracy to improve, effort to fall, or reuse to generalise. Several exit criteria are written in the form "measured, **or** a recorded and reasoned statement of why it is not measurable", and both branches satisfy the plan. This is not defensive drafting: it follows from the fact that whether governed, reused expert judgment improves agentic assessment is an *empirical outcome* of the research, and a plan whose milestones depend on the answer coming out one way would create pressure to produce that answer.

6.3 Dependency view

6.3.1 Gate register

Gate	Description	Current state	Owner
G-RQ	D-RQ-01 and D-RQ-02: umbrella RQ and SQ1–SQ3 wording signed off	Pending; wording provisional from the 2026-08-05 call, unverified against the candidate's saved draft	Iris and Arnon
G-LIT	QL-01 to QL-05 executed, screened, deduplicated, and synthesised	Not run. All query and database rows are <i>Protocol ready / not run</i>	Ali
G-LBL	At least 20 adjudicated generalization-safe EXP-005 labels	0 of 24 supplied. 0 of 2 independent reviewer returns	Reviewers, adjudicator, supervisors
G-UNI	A-14: written confirmation of the university proposal and candidacy process	Open; all eleven fields Pending; owner unassigned	Unassigned — assignment is itself overdue
G-MED	Medical entry gates G1–G6 (use case, people, authorisation, ethics/privacy, environment, protocol)	0 of 6 passed. Verdict NO-GO/BLOCKED	Supervisors, partner, custodian, ethics owner, IT/security
G-PROTO	Supervisor acceptance of each study protocol before that study executes	Recurring, per study	Iris and Arnon
G-2ND	Written authorisation for the second setting used in Study 3	Not yet selected	Ali and supervisors

6.3.2 Block-to-gate dependencies

Block	Cannot start until	If the gate stays closed
Block 0	—	—
Y1 B1 (searches)	G-RQ	Searches keyed to unconfirmed concepts would have to be re-run; the block converts to protocol refinement and the review slips one block
Y1 B1 (Chapter 2)	G-LIT	No novelty or review-completeness claim may appear anywhere in the proposal or thesis; Chapter 3's gap statements remain candidate claims
Y1 B1 (Study 1 freeze)	G-RQ, G-PROTO	Construction cannot begin; Y1 B2 slips
Y1 B2 (gold set)	G-LBL	No accuracy, macro-F1, effort-reduction, superiority, or generalization claim is available anywhere in the thesis. Studies 1 and 3 report readiness and mechanism evidence only. This is the

		single most consequential dependency in the plan
Y1 B3 (Study 1 results)	G-LBL	Study 1 reports the inspectable architecture, the analytical burden model, scenario validation, and the preregistered empirical test, with effectiveness claims excluded
Y1 B4 (Study 2 freeze)	G-PROTO	Y2 B1 construction slips
Y2 B2 (Study 2 evaluation)	Study 1 and 2 artifacts; recruited expert panel	Mechanism and contract validation only; safe reuse is not inferred from fixtures
Y2 B3 (Study 3 preregistration)	Studies 1–2 artifacts; G-PROTO	The transfer arm cannot open; opening it without preregistration would void the transfer evidence
Y2 B4 (transfer arm)	G-2ND (Plan B) or G-MED all six (Plan A)	Study 3 reduces to within-corpus transfer across settings and diagram types; the cross-institution arm is reported as blocked with its cause
Y3 B1–B2	Y2 B4 evidence	Transfer classification is reported as untested rather than assumed
Y3 B3–B4 (submission)	G-UNI	Submission dates remain internal targets; the candidate cannot know the deadline, reviewer count, or required lead time. This gate does not block <i>research</i> , only <i>scheduling</i> — which is precisely why the plan does not let it block anything else
Any Plan A activity	G-MED , all six, cumulatively and sequentially	No row-level work, pilot, export, or medical claim. Plan B proceeds unchanged

6.3.3 Two structural properties of this dependency graph

No research question depends on a gate outside the candidate's and supervisors' reach. G-MED is held by external institutions; it gates only the optional Plan A extension. G-UNI is held by the university; it gates only dates, not content. Everything that determines whether the doctorate is answerable — G-RQ, G-LIT, G-LBL, G-PROTO, G-2ND — is held by the candidate and the supervision team.

The expert-label gate is the critical path. G-LBL sits upstream of every quantitative statement in the thesis. It is currently at zero. It is also the gate with the shortest path to clearing: it needs two named reviewers, a signed approval pack, and adjudication — no institution, no ethics board, no partner. This is why Block 0 and Y1 B1 both carry labelling mobilisation as a named deliverable, and why Y1 B2's exit criterion explicitly permits a recorded failure rather than pretending the gate will clear.

6.4 The Plan A / Plan B fork

6.4.1 The invariant core

Both plans carry the **same umbrella research question and the same SQ1–SQ3**. Both complete Study 1 (intervention architecture) and Study 2 (governed judgment lifecycle) in the software/modelling context. Both complete Study 3's evaluation in software/modelling through independent human evidence and explicit claim gates. Both include the full literature synthesis, the domain-neutral artifact and governance model, and a documented second-setting transfer comparison.

The difference between the plans is confined to **which second setting Study 3 uses**, and nothing else.

6.4.2 The same twelve blocks under Plan B

Under Plan B — no medical access at all, at any point in three years — the block grid in §6.2 is executed **unchanged**. Only three cells differ, and all three sit inside Study 3:

Block	Plan-B content	Difference from §6.2
Y1 B4	Medical gate re-review still logged (as a recorded NO-GO)	None to the research; the gate review remains as a governance record
Y2 B3	Second setting selected from the authorised non-medical candidates: a new diagram family or modelling task, a second annotated corpus or run, an external course or institution, a different reviewer panel or expertise profile, or a longitudinal reuse setting across iterations	Selection source only
Y2 B4 – Y3 B1	Transfer arm executes across software/modelling settings, institutions, datasets, diagram types, or reviewer panels	Setting only; design, blinding, adjudication, leakage controls, and transfer classification are identical
Y3 B2	Paper 3B — cross-domain / cross-institution software/modelling evaluation-and-transfer paper	Venue and framing only

Read against §6.2, this table is the substantive claim of the section: **every research question is answered, every study is completed, every paper is written, and the thesis is submitted, with no medical access whatsoever**. Plan B is the protected default, not a degraded fallback. It preserves construct comparability, independent human evidence, leakage control, and the same SQ3 evaluation and transfer measures.

6.4.3 What Plan A would add

Plan A is **additive only**. It never substitutes for Plan B work and never displaces a Plan B milestone. If all six gates clear, the following work lands inside the existing blocks:

Added work	Lands in	Prerequisite
Clinical use case defined: workflow, decision point, problem owner, unit of analysis, input/output contract, baseline, non-goals,	Y1 B4 or later, whenever G1 clears	G1

measurable success and failure criteria		
Named clinician, data custodian, privacy/ethics owner, VDI administrator, methods reviewer, with availability and escalation recorded	Same block as G1	G2
Individual project-specific authorisation for each researcher, with named data source, least privilege, and expiry	Following block	G3
Written ethics and privacy determination covering the selected data, derivatives, linkage, retention, publication, disclosure, and incidents	Following block	G4
Approved restricted environment: VDI, storage, compute, audit logging, egress controls, allowlisted offline/no-telemetry tooling, and either an approved local model or an explicit <i>No LLM</i> decision	Following block	G5
Approved, versioned clinical protocol: cohort, inclusion/exclusion, windows, outcome, comparator, missingness, leakage controls, statistics, and stop rules	Y2 B3, alongside Study 3 preregistration	G6
Downstream controls D1 (data integrity and provenance), D2 (bounded pilot), D3 (disclosure and export) — each separately gated, none authorised by passing G6	Y2 B4 onwards	All six gates, sequentially
Medical workflow requirements used as a design stress test for Study 1	Retrospective, after expert review	G1, G2
Paper 3A replaces Paper 3B	Y3 B2	Complete gated evidence

The gates are **sequential and cumulative**: use case → people → authorisation → ethics/privacy → environment → protocol. A gate passes only when every required item is verified, current, project-specific, and approved by every named approver; partial evidence produces no partial pass. If any of the use case, people, data source, authorisation, determination, environment, or protocol changes materially — or an approval expires — that gate and every later control return to blocked.

Because each addition sits *behind* a gate held by an external institution, none of them can be dated. In the vocabulary of §6.1.3 they are **Contingent**, and the block positions above are provisional placements, not schedule commitments.

6.4.4 The 2026-08-26 checkpoint and the reopening rule

The fork is decided by an internal control date: **26 August 2026, end of day, Asia/Jerusalem**. The decision rule is objective — if any of G1–G6 lacks a documented owner, an evidence path, and a feasible completion date at that checkpoint, Plan B becomes the committed execution path from the following day, and Plan A is written into the proposal as a contingent extension.

Three qualifications matter and should not be dropped in revision:

22. This date is an **internal project-control date**. It is not a university deadline, and it is not attributable to any supervision meeting. It exists so that the proposal converges on an executable path rather than staying ambiguous.
23. The current state is **0 of 6**, with a further unresolved complication: Gates 1 and 3 each require a *named* data source, and a third dataset (EHRSHOT) observed on 2026-08-11 is named in none of the existing gate documents, all of which were written against MIMIC and Clalit. On present evidence the checkpoint resolves to Plan B.
24. **Reopening is possible but controlled**. After fallback, Plan A can reopen only through completed G1–G6 evidence, the applicable downstream controls, an explicit supervisor decision, governance approval, and a schedule-impact review — and it cannot displace committed Plan B milestones without a recorded change. In block terms: Plan A can be added to a later block; it cannot reclaim an earlier one.

6.5 Milestones and review points

6.5.1 Weekly supervision cadence

Status: Verified. A recurring supervision series is accepted by Ali, Iris, and Arnon — Wednesday 09:00–10:00 Asia/Jerusalem — with acceptance confirmed through the 2026-10-07 occurrence. Continuation beyond that occurrence is expected but is not yet verified and should not be asserted as such.

The weekly meeting is the plan's week-level control instrument, which is what allows the written plan to remain at block resolution. Each meeting follows a fixed shape: a short pre-read circulated in advance; completed artifacts shown with evidence links; blockers and required decisions stated explicitly; every decision recorded as approved, corrected, deferred, or rejected; one primary next task assigned with an owner and an acceptance check; and the proposal, action register, risk register, and literature workbook updated afterwards.

The cadence verifies logistics only. It approves no research decision by itself; approval exists only where a decision is logged.

6.5.2 Block-boundary reviews

Status: Structural. Four per year, at each block boundary, twelve across the plan. Each is a longer session with a fixed agenda:

25. Test the closing block's exit criteria one by one and record pass, partial, or fail with evidence.
26. Re-test each open gate in §6.3.1 and log its state.
27. Confirm or re-baseline the next block's deliverables.
28. Re-review the claim boundary: is any claim now being made that its evidence does not support, and is any claim now supportable that is still being withheld?
29. Record every change as a decision with old wording, new wording, rationale, approver, and effective date.

The fourth item runs in both directions deliberately. Over-claiming is the more serious failure, but under-claiming — leaving evidence unreported because a boundary was set conservatively at an earlier stage and never revisited — is also a failure, and a block-boundary review is the right place to catch it.

6.5.3 Candidacy and submission milestones

Status: Internal working target — every item below is unverified.

Milestone	Working position	Verification required
Proposal submission	Block 0 / early Y1 B1	Official deadline, forms, submission channel, length and formatting rules, language, reference requirements
Reviewer appointment	Unknown	Reviewer count and eligibility; nomination and conflict-of-interest process; committee composition
Candidacy presentation	Unknown	Whether required at all; duration, audience, assessment criteria; required lead time from submission
Annual progress reporting	Y1 B4, Y2 B4	Whether required, in what form, assessed by whom
Thesis submission	Y3 B4	Process, forms, examiner appointment, lead times

The University Process Verification Checklist records all of these as *Pending*, with the owner unassigned. The candidacy presentation should be prepared only after these requirements are verified — preparing it against assumed rules risks preparing the wrong artifact.

Assigning an owner to A-14 and authorising the inquiry is the single cheapest item in this entire plan and it is currently the only one blocking an entire class of milestones. It carries no research risk and no evidence risk; it requires one email to an authoritative office.

6.5.4 Publication milestones

Status: Internal working target. Three submissions, one per study: Paper 1 in Y1 B3, Paper 2 in Y2 B3, Paper 3 (3A or 3B) in Y3 B2. **Submission**, not acceptance, is the milestone throughout. Acceptance timing is controlled by reviewers and venues and is therefore never used as a gate for any downstream block; a rejected paper is re-targeted within the following block without disturbing the grid.

6.6 Risks to the schedule

#	Risk	Trigger	Schedule effect	Mitigation	Owner
R1	Expert labels never arrive	Reviewers unavailable, unwilling, or unable to complete blinded labelling and adjudication	Y1 B2 and Y1 B3 lose their quantitative content; the thesis becomes readiness- and mechanism-evidence only	Recruit reviewers in Block 0, not later; keep the labelling instrument short and resumable; recruit a third reviewer as substitute; accept and report the shortfall explicitly rather than substituting synthetic or model-generated labels under any	Ali + supervisors

				circumstances	
R2	RQ wording sign-off slips	D-RQ-01/D-RQ-02 remain pending past Block 0	Literature searches cannot execute against confirmed concepts; Y1 B1 slips one block and the effect cascades	Bring a decision-ready read-back with a recommended wording and two variants to a weekly meeting; treat "no decision" as itself a logged decision with a schedule consequence	Iris and Arnon
R3	University process never verified	A-14 stays unassigned or unanswered	Submission and candidacy cannot be scheduled; the risk is discovering a passed deadline retrospectively	Assign an owner at the next weekly meeting; escalate to the department if unanswered within one block; keep every date labelled as an internal target until answered so no false commitment is created	Ali + supervisors
R4	Medical gates absorb effort without clearing	Plan A work continues past the 2026-08-26 checkpoint without gate progress	Direct loss of blocks from the protected core	Objective checkpoint rule; Plan B is the committed default; a strict stop rule on medical familiarisation; the reopening rule prevents Plan A from reclaiming committed blocks	Ali + supervisors
R5	Literature review becomes descriptive	Paper-by-paper summaries accumulate without synthesis	Chapter 2 slips and, worse, the gap argument stays unsupported	One paper per row with mandatory synthesis fields (gap relevance, SQ/study mapping, contradiction, contribution use); authors' conclusions kept separate from the candidate's interpretation; a synthesis memo at every block boundary	Ali
R6	Second setting cannot be authorised	No external course, institution, corpus, or reviewer panel agrees in time	Y2 B4 transfer arm cannot open; Study 3 reduces to within-corpus transfer	Identify and approach candidate settings a full year ahead, in Y1 B4; maintain at least two live candidates; within-corpus transfer across diagram types and domains is a legitimate reduced arm and is reported as such	Ali + supervisors
R7	Evidence overclaim	Quantitative or medical wording appears without eligible evidence	Not a schedule risk but a credibility risk that can force retraction and rework across chapters	The 0-of-24 and 0-of-6 boundaries restated wherever the evidence is summarised; a claim-and-evidence review at every block boundary; explicit "not yet computable" language	Ali

R8	Artifact construction expands to fill the blocks	Design-science construction continues past its exit criterion because there is always more to build	Y2 and Y3 evaluation blocks are compressed	The exit criteria in §6.2 are coverage criteria (every requirement maps to an inspectable artifact and a validation), not completeness-of-implementation criteria; construction stops when coverage is reached	Ali + supervisors
R9	Paper rejections cascade	Any of Papers 1–3 is rejected	None, by construction	Submission rather than acceptance is the milestone; re-targeting happens inside the following block; no downstream block depends on an acceptance	Ali
R10	Single-person dependency	Illness, reserve duty, employment demands, or any other extended unavailability	Direct block-for-block slippage	Y3 B4 is the structural buffer; block-level rather than month-level granularity absorbs short absences without re-planning; artifacts and decisions are recorded continuously so work is resumable by the candidate after an interruption	Ali + supervisors

6.6.1 The two risks that actually govern this schedule

Of the ten, two dominate.

R1 (expert labels) is the critical path, for the reason given in §6.3.3: it sits upstream of every quantitative claim, it is currently at zero, and yet it is entirely within the supervision team's reach. Its mitigation is not a process control but a decision — name the reviewers and release the approval pack. Everything else in the mitigation column is secondary to that.

R10 (single-person dependency) is the risk the plan structurally cannot mitigate away, because a doctorate is by definition the work of one person. The block structure is the mitigation: a three-month window absorbs a two-week absence without any re-planning at all, whereas a month-level plan would require re-baselining. This is the last and most practical argument for the granularity that E14 requires — a plan expressed in blocks is a plan that survives contact with a real year.

Chapter 6 prepared as a complete working draft. Research-question wording follows Chapter 3 and remains provisional pending logged D-RQ-01 and D-RQ-02 decisions. All references to the 2026-08-05 supervision discussion derive from a machine transcript with inferred, undiarized speakers; attributions are the record's inferences, pending the participants' own confirmation. EXP-005 stands at 0 of 24 generalization-safe labels supplied against a threshold of 20. The medical track stands at 0 of 6 mandatory entry gates. Literature queries QL-01 to QL-05 are frozen and not run. No date in this chapter is a confirmed university deadline.