# Sections 2 and 4 — Thinking Notes

## Sections 2 and 4 — Thinking Notes (options only, nothing started)

> **Status: THINK-ONLY. Deliverable `A08-04` from the 2026-08-05 call.**
>
> Iris's instruction was explicit: *think about* sections 2 (literature survey) and 4 (research artifact per research question), but **do not start them**. This note therefore contains **options with trade-offs and open questions** — no committed chapter design, no chosen artifact, and no executed literature search.
>
> **Boundary facts carried into every option below.** The searches `QL-01`–`QL-05` are protocol-ready and **not run**, so no option may assert review completeness or novelty. EXP-005 holds **0 of 24** required independent generalization-safe expert labels, and **0 of 6** medical entry gates pass, so no option may assume accuracy, generalization, clinical-performance or effort-reduction evidence. The RQ/SQ wording is **provisional** pending `A08-01` verification and logged `D-RQ-01`/`D-RQ-02` decisions — which is itself a live design constraint on §2's structure.

---

### Part 1 — §2 Literature survey: four structural options

Four genuinely different ways to organise the chapter. **None of these is a recommendation.** The choice interacts with the wording sign-off: options that hard-code sub-question wording into headings become expensive if the wording changes on Aug 12.

#### A. SQ-parallel chapter (one section per sub-question)

**Structure.** 2.0 scope + review-method statement (protocol described, results NOT claimed, since QL-01..QL-05 are PROTOCOL READY / NOT RUN) -> 2.1 literature bearing on SQ1 (selective intervention) -> 2.2 SQ2 (governed knowledge reuse) -> 2.3 SQ3 (evaluation and transfer) -> 2.4 synthesis handing three candidate gap statements to Chapter 3. Mirrors the SQ=Study 1:1 mapping in three-study-contract.md.

**Pros.** Mechanically inherits the tagging that already exists: literature/per-rq-literature-map.md already tags sources RQ1/RQ2/RQ3/general, and literature-search-execution-register.md already maps each frozen query to a primary SQ (QL-01->SQ1, QL-02->SQ2, QL-03->baseline/Studies 1-3, QL-04->SQ1+SQ3, QL-05->conditional Plan A only). Chapter 2 -> Chapter 3 handoff is trivial because chapter-3-gap-and-research-questions-draft.md 3.1 already argues exactly three gap pillars in the same order. Each section is separately extractable as the related-work section of Paper 1 / 2 / 3.

**Cons.** Hard-codes PROVISIONAL SQ wording (D-RQ-01/D-RQ-02 unsigned) into chapter headings, so a wording correction on Aug 12 reshapes the chapter's skeleton. Literature does not partition by our questions: the sources currently tagged 'general' in the per-RQ map (HAI-001, GOV-001, MDE-001 by manifest ID) must be either duplicated across sections or parked in a residual bucket. It structurally exposes the thinnest area - the per-RQ map records RQ3 as having no tagged sources yet - as a visibly empty section. Highest risk of reading as an advocacy chapter (each section built to justify its own SQ) rather than a survey.

**Fit to the three-question spine.** Highest structural fit to the 3-SQ spine, but the fit is conditional on wording sign-off. Strong candidate IF D-RQ-01/D-RQ-02 are signed on Aug 12; weakest of the four if the wording stays open. Not a recommendation - an option.

## B. Community/discipline-parallel chapter (follow the existing taxonomy branches)

**Structure.** Sections follow docs/research/literature-review-taxonomy.md as it already stands: the eight original branches (human-in-the-loop, human-on-the-loop, explainable AI, expert feedback and knowledge capture, AI-assisted domain modeling, model assessment and variability interpretation, human-AI co-reasoning, design science), then the six July-2026 redirect branches (agentic HITL architectures, human-AI collaboration in multi-agent systems, RLHF/feedback-learning contrast, agent memory and learning from feedback, anti-sycophancy/source-grounded challenge, intervention policies and workload/dosage, evaluation of human-AI systems), with the three MediVARIA branches at explicitly lower priority.

**Pros.** Matches how the reading is already organized, so no re-tagging cost. Fully decoupled from the SQ wording - survives any D-RQ-01/D-RQ-02 correction unchanged. The taxonomy file already scopes its July-2026 section as the survey scope for Pnina's research-methodology course, so a community-shaped chapter serves the course deliverable and the proposal from one structure. Makes it natural to state what each community does and does not address without asserting our gap as established fact - which is required while the searches are unrun.

**Cons.** Eight + six + three branches is far too many sections for a proposal chapter; a merge is required and the merge decisions are unresolved. The gap argument becomes implicit and has to be reassembled in a synthesis section anyway, so the SQ payload is deferred rather than avoided. High duplication risk - several candidate works would sit naturally in two or three branches. Reviewers may not see how the survey drives the three questions.

**Fit to the three-question spine.** Moderate fit to the 3-SQ spine: it serves the reading work and the course deliverable well but leaves the entire SQ mapping to a synthesis layer that then does the real work. Best option if the chapter must be robust to a wording change on Aug 12.

## C. Lifecycle-stage chapter (organize by stages of the judgment lifecycle)

**Structure.** Sections follow the 'Review Lens' sentence already in the taxonomy file (judgment selectively triggered, structurally captured, stored as reusable knowledge): (a) deciding to ask / trigger and selective intervention; (b) capturing judgment and the system's own reasoning; (c) validating and challenging (source-grounded verification, anti-sycophancy); (d) reconciling, authority, adjudication, revocation; (e) storing and scoping (provenance, memory, scope); (f) retrieving and reusing (explainable retrieval, advisory); (g) evaluating and transferring. Design science and governance/risk run as cross-cutting subsections.

**Pros.** Mirrors the thesis's own claim shape - the taxonomy's 'Thesis Use' section already frames the novelty as the combined lifecycle (selective trigger, structured capture, transparent reusable memory, controlled reuse), so the chapter's spine and the contribution's spine are the same object. Naturally surfaces the contrast the taxonomy's reading rules explicitly ask us to record: work that handles one-time feedback and stops. Stage boundaries are more stable than SQ wording.

**Cons.** The lifecycle is OUR construct. Organizing the field by our own artifact's stages is exactly the solution-first framing flagged in the 2026-08-05 record (item E4, solution-vs-question conflation) and invites the criticism that the literature was fitted to the design. Whole communities (evaluation of human-AI systems, design science) do not sit at any single stage. The conditional medical branch (QL-05) has no natural home.

**Fit to the three-question spine.** Good fit to SQ CONTENT but structurally unbalanced against the 3-SQ spine: SQ1 maps to one stage, SQ2 to roughly five, SQ3 to one - so SQ2 would dominate the chapter while Study 1 and Study 3 look under-reviewed. Carries the highest E4 risk of the four.

### D. Hybrid two-layer: compressed community body + design-space synthesis resolving to the SQ gaps

**Structure.** Layer 1 (body): the community sections of option B merged down to roughly five. Layer 2 (synthesis): position the reviewed work along a small set of explicit design-space axes drawn from contrasts Chapter 3 3.1 already argues - blanket vs selective oversight under an attention budget; static knowledge base vs weight-absorbed feedback vs a governed middle; scope-carrying vs unscoped reuse; single-context vs cross-context evaluation - presented as a work-by-axis matrix, closing with the three candidate gaps handed to Chapter 3.

**Pros.** Keeps a descriptive survey (course deliverable, robust to a wording change) while still delivering the SQ-shaped payload. The axis matrix makes 'what is missing' visible as a position map rather than as an assertion of completeness. Positioning rather than listing is what the taxonomy's reading rules ask for ('capture both mechanism and evaluation'). The matrix is also the honest place to record current coverage state - e.g. that RQ3 has no tagged sources yet - without dressing it up.

**Cons.** The most work: two organizing schemes must be built and kept mutually consistent, and it produces the longest chapter. The axes are themselves a design decision that needs supervisor buy-in before writing, so it front-loads a discussion. Real hazard: a filled matrix reads as a completeness claim unless it is explicitly labelled as coverage of what was screened - and with QL-01..QL-05 unrun, no completeness or novelty statement is permitted at all.

**Fit to the three-question spine.** Best fit to the 3-SQ spine without hard-coding provisional wording into headings, at the highest cost. Only viable if the searches actually execute with enough lead time to populate the matrix; otherwise the matrix would be mostly empty and would misrepresent the evidence state.

## Part 2 — §4 Research artifact per research question: options

For each sub-question, the option the three-study contract already names ("contract-anchored") plus genuine alternatives at different abstraction levels. Listed to expose the **granularity and abstraction decisions** that have not yet been made — not to select one.

A recurring issue visible across all three: the contract's current "Core artifact" rows bundle many components into one artifact (six for SQ1, nine for SQ2, ten for SQ3), and several of the SQ3 items are *outputs of running the study* rather than the artifact itself. Whether a study may have a **package** or must have exactly **one named artifact** is an open question for the supervisors.

### SQ1 — Selective intervention

| Option | What it would be | Evaluation shape | Risk | Hard dependency |
|---|---|---|---|---|
| **A1. Contract-anchored bundled selective-intervention reference architecture** | The Study 1 'Core artifact' row taken as written - listener/event catalog, eligibility/priority and uncertainty criteria, triage/dosage policy, human-routing contract, timeout/escalation rules, burden budget - packaged as one domain-neutral reference architecture, with the existing VEGO-AI selective-intervention policy and review queue as a single instantiation. Sharpening the contract row does not yet do: that row bundles at least three artifact TYPES (a structure, a policy, and a budget model). The option is only coherent if one is designated the contribution and the others are scoped as its supporting mechanism. | Requirements-to-artifact coverage (every SQ1 intervention requirement traces to a component and a validation), architecture traceability, bounded scenario analysis - all as listed in the contract's 'Primary method' row - plus the already-existing offline replay series (EXP-006 event reconstruction, EXP-007 per-mode load-vs-coverage, EXP-008 unstable-but-never-reviewed trigger candidates) reported strictly under their stated claim scope: observability / mechanism / design evidence only. | Reads as an engineering deliverable rather than a doctoral contribution unless an explicit generalizable claim sits on top - the same objection the 2026-08-10 brainstorm raised against its own architecture option. Bundling six components hides which one is the claim. Sharpest claim hazard: EXP-007's 'load reduction vs every_decision' is a routing-count property of an offline replay and must never be phrased as expert-effort reduction. | Supervisor decisions on eligibility, dosage, timeout, and routing (named as Study 1 key dependencies); a stable event/uncertainty contract; D-RQ-01/D-RQ-02 wording sign-off; QL-01/QL-04 execution to ground the requirements' provenance (both PROTOCOL READY / NOT RUN). No EXP-005 dependency, which is why SQ1 is the least gate-blocked of the three. |
| **A2. Attention-budget cost/coverage model as the artifact (instrument, not architecture)** | Rather than an architecture, the artifact is an explicit analytical model relating trigger configuration to (i) expected reviewer load and (ii) coverage | Analytical evaluation of model properties (monotonicity, boundary behaviour, degenerate cases), instantiated by offline replay on the existing four settings using EXP-007's per-mode | A model whose parameters are only ever fitted to one corpus invites the question 'is this a model or a fit to VEGO-AI?'. Second risk is direction-of- | No expert labels are needed for the mechanism layer, which is this option's main attraction. But any statement that a given budget setting preserves |

| | | | | |
|---|---|---|---|---|
| | of uncertainty-marked events, parameterised so a site sets a budget and reads off an admissible policy region. The four dosage modes already replayed (every_decision / threshold / first_n_then_auto / silent) become four points in the model's space rather than the artifact itself. Distinct from the brainstorm's 'decision framework' option: the contribution is the cost/coverage relation, not a policy table. | routed-item counts and uncertainty coverage with stable denominators, plus sensitivity/scenario analysis. Explicitly excludes any quality, accuracy, or effort claim - the contract's 'Measures explicitly excluded until gated' row forbids optimal dosage and workload-reduction claims based only on architecture/tests/fixtures. | travel: EXP-007 was specified to report a load-vs-coverage frontier and explicitly NOT to select or silently tune a default - an artifact framing could quietly convert that frontier into a recommended operating point, which the experiment plan forbids. | assessment quality is blocked until EXP-005 has at least 20 generalization-safe adjudicated labels; it currently has 0 of 24. |
| **A3. Stated design principles with nested instantiations (higher-abstraction option)** | The artifact is a small set of explicitly stated, testable design principles for selective intervention in agentic assessment under an attention budget - each written as condition -> design move -> expected consequence -> falsifier - with the VEGO-AI implementation demonstrated as one instantiation and at least one alternative instantiation sketched. This answers the 'is a component diagram doctoral?' objection at the level of abstraction rather than by adding more components. | Principle-level evaluation: supervisor/expert walkthrough for clarity and non-triviality; demonstration that each principle is instantiable in at least two distinct designs; and a mapping from each principle to at least one falsifiable measurement that a later, gated study could actually run. No empirical principle validation at proposal stage. | Principles not derived from a completed review look asserted rather than grounded - and the review that would ground them has not been run, so every principle is a candidate claim, not a finding. Second risk: at this abstraction the artifact can drift into restating the SQ. Methodological framing for a 'design principles / design theory' contribution type: GAP - no verified source in the repo; the taxonomy has a design-science branch but no verified citation behind it. | QL-01/QL-04 execution for principle provenance (not run); supervisor agreement that a principle set rather than a system counts as the Study 1 artifact; D-RQ-01/D-RQ-02 sign-off, since the principles would be phrased against the SQ1 wording. |

## SQ2 — Governed knowledge reuse

| Option | What it would be | Evaluation shape | Risk | Hard dependency |
|---|---|---|---|---|
| **B1. Contract-anchored governed** | The Study 2 'Core artifact' row taken as written: structured judgment | Entirely artifact-side and available without real judgment records: | An 'everything' artifact whose boundary reviewers will contest. | M-04 protocol approval before EXP-009/010 |

| | | | | |
|---|---|---|---|---|
| **judgment lifecycle (bundled, as written)** | schema, H-Verify/source challenge, conflict/adjudication and authority rules, provenance model, judgment memory, explainable retrieval, scope controls, expiry/revocation, and bounded integration/reuse. Sharpening: that is nine named components. As with SQ1, the option is viable only if one component is designated the contribution and the remaining eight are scoped as its supporting mechanism - otherwise the 'framework/protocol/process' boundary objection from the 2026-08-10 brainstorm lands. | schema/contract validity, provenance completeness, validation and reconciliation coverage, conflict/authority rule coverage, retrieval-explanation completeness, unsafe-reuse rejection, scope/expiry/revocation enforcement, bounded convergence - the contract's own SQ2 primary-measures list. | Concrete evidence hazard: the H-Verify and convergence halves currently rest on EXP-009/EXP-010, which the experiment plan labels PROVISIONAL synthetic fixtures, assumption-driven, isolated as SYNTHETIC_NOT_HUMAN, and gated on M-04 protocol approval before results may be interpreted or rerun. They cannot be presented as validation of real expert-error handling. | results are interpreted; Study 1's intervention points; supervisor decisions on validation, conflict, authority, convergence, and reuse scope. Anything beyond mechanism/contract validation needs real approved judgment records - none exist; the contract's own fallback says do not infer safe reuse from fixtures. |
| **B2. Normative judgment-record contract plus executable conformance suite (specification-as-artifact)** | A system-independent normative specification of what a reusable judgment record must carry - case grounding, the system's core reasoning as the expert actually saw it, the expert's rationale, scope, authority, provenance chain, validity/expiry - paired with an executable conformance suite any implementation can be run against. VEGO-AI becomes reference implementation #1; the contribution is the contract plus the test that decides conformance. This is deliberately NOT the brainstorm's 'schema architecture' option, which describes our schemas; this is a normative spec that an implementation can fail. | Conformance evaluation: our implementation passes; at least one deliberately non-conforming variant is shown to fail for the right, named reason; and the spec itself is inspected for completeness and ambiguity. The existing architecture-conformance series (EXP-013..EXP-018: schema, lineage, explicit-gap, determinism via identical normalized hashes, isolation and non-application checks) is the shape this would generalize - currently offline fixture evidence granting no runtime authority. | A specification with exactly one implementation is a schema with a longer name; the doctoral claim rests on demonstrating implementation-independence, which needs a second implementer. Equally important: conformance is not safety - passing the suite says nothing about whether any reuse is actually safe, and that gap must be stated in the artifact, not discovered by a reviewer. | An independent implementer or reviewer to test implementation-independence: GAP - no named party in the repo. Supervisor agreement that a spec plus conformance suite is an acceptable artifact class. No EXP-005 dependency, which makes it buildable now. |
| **B3. Authority-and-revocation model, evaluated against an unsafe-reuse failure-mode catalogue** | Carve out the half of SQ2 that the question wording itself names as a failure mode - loss of human authority - and make that the artifact: who may assert, endorse, contest, expire, and revoke a judgment; what happens to downstream reuse when authority is withdrawn; how an expert can see where their judgment travelled. | Failure-mode coverage (each catalogued mode has a control and a test showing the control fires), revocation-propagation tests over existing memory artifacts, authority-transition trace inspection, and a walkthrough with an accountable expert on whether the model | Narrower than the contract row - supervisors may read it as one component of Study 2 rather than as Study 2. And the catalogue's completeness cannot be claimed: it would be an enumeration from our own design, not from an executed review, so it is a candidate taxonomy | QL-02 execution to ground the catalogue in prior failure taxonomies rather than introspection (PROTOCOL READY / NOT RUN). At least one accountable expert for the walkthrough - available in the software/modeling |

| | | | | |
|---|---|---|---|---|
| | Paired instrument: an enumerated catalogue of unsafe-generalization and authority-loss failure modes, each with a detection control. Distinct from the brainstorm's 'safe-reuse classification scheme', which classifies a judgment at match time; this governs the human's standing over the judgment's whole life. | matches how they expect to retain control. | only. | instantiation; for the medical instantiation this is gate-blocked (named expert with committed availability is Plan A readiness item A2, unevidenced; medical entry gates are 0/6). |

## SQ3 — Evaluation and transfer

| Option | What it would be | Evaluation shape | Risk | Hard dependency |
|---|---|---|---|---|
| **C1. Contract-anchored evaluation-and-transfer package (as written)** | The Study 3 'Core artifact' row as it currently stands: preregistered evaluation and workload protocol; adjudicated reference evidence; result/validity package; traceability audit; transfer taxonomy; common-core/adaptation map; domain contract; governance profile; replication report; stop/fallback rules. Sharpening the row needs: at least six of those ten are OUTPUTS of running the study, not artifacts designed by it. The honest version of this option declares the protocol, the domain contract, and the stop/fallback rules to be the artifact, and reclassifies the rest as results. | Pre-execution, protocol-level evaluation only: preregistration review, leakage-control inspection, sealed development/holdout discipline, no-baseline-overwrite check, and a documented stopping rule. Everything execution-side (paired comparison, inter-rater agreement, adjudication and escalation rates, net correction) is entirely post-gate. | A protocol reads as methodology rather than artifact, which is precisely the solution-vs-question conflation flagged on 2026-08-05 (item E4) and already noted as a con in the 2026-08-10 brainstorm. Larger hazard: while EXP-005 is 0/24 the package's evidential half cannot be produced at all, so presenting the ten-item package as 'the artifact' risks implying an evaluation exists. | EXP-005 at 20+ generalization-safe adjudicated labels (currently 0 of 24, from 27 blind rows); two independent reviewers plus an adjudicator (unnamed); frozen baseline with no overwrite. Plan A additionally requires all six entry gates G1-G6 to have owner, evidence path, and feasible date by the internal 2026-08-26 checkpoint - currently 0/6, so Plan B activates by design if any gate fails. |
| **C2. Transfer-eligibility decision procedure plus target-context descriptor** | An operational procedure taking (judgment record, target-context descriptor) and returning eligible / eligible-with-stated-adaptation / blocked, with a recorded reason - plus the schema of the | Decidability and reliability OF THE PROCEDURE: two trained raters independently apply it to the same set of existing judgment records and context descriptors; report agreement, undecidable cases, | Agreement on APPLYING a procedure is not evidence that its verdicts are correct. If reported without that framing it will be misread as transfer evidence, which is a forbidden claim while EXP-005 is 0/24. | A supervisor ruling on whether instrument-reliability evidence is admissible ahead of the EXP-005 gate; two raters (non-clinical raters suffice for the software/modeling corpus); existing judgment records to |

| | | | | |
|---|---|---|---|---|
| | context descriptor itself (the contract's 'domain contract' pulled forward as the designed part). This sharpens the brainstorm's 'transfer classification scheme' from a taxonomy into a decidable procedure, and it sharpens the SQ2/SQ3 boundary the brainstorm flagged as unresolved: SQ2 records reuse_scope ON the judgment; SQ3 decides eligibility AGAINST a described target. | and the distribution of blocking reasons. This evaluates the instrument, not assessment accuracy - so it is one of the few Study 3 evaluation paths with no dependence on expert gold labels, subject to supervisors confirming that instrument-reliability evidence is admissible. | Secondary: it still needs two raters, so it is gold-label-free, not label-free - do not oversell it as unblocked. | apply it to. Any healthcare application remains blocked at 0/6 medical gates. |
| **C3. Cross-context construct-comparability instrument** | The artifact is the measurement-validity layer: an explicit mapping of the constructs any transfer claim depends on - what counts as a deviation, what makes an uncertainty important, who qualifies as the expert, what a 'context' is - between the software/modeling baseline and any second context, so that 'the same measure' provably means the same thing on both sides. The Plan B row already lists a construct-comparability and provenance record as mandatory readiness evidence; this option promotes that compliance item to the designed artifact. | Construct-mapping inspection in which explicitly non-mappable constructs are reported as a result rather than a failure (a documented 'does not transfer as a construct' list is a finding); expert review of the mapping; and a demonstration on the only comparison currently stageable - within software/modeling, across the settings and diagram types already in the existing corpus. | The option most likely to be judged infrastructure rather than contribution. And its most interesting instance - software/modeling versus healthcare - cannot be built now: no medical use case, dataset, expert, or approval is evidenced, Plan A items A1-A8 are unmet, and MIMIC is recorded as an exploratory resource that is not the selected dataset. | An authorized second context. Plan B needs an authorized independent software/modeling corpus plus a reviewer plan - not yet named. Plan A needs all six entry gates (0/6) by the internal 2026-08-26 checkpoint. Neither is satisfied today, so at proposal stage this artifact can only be specified and demonstrated within-corpus. |

## Part 3 — Open questions for Iris and Arnon (Aug 12)

These are the decisions that must be made before §2 or §4 can legitimately start. They are grouped so they can be worked through quickly in the meeting.

**Chapter 2 shape and scope**

1. Chapter 2 shape: which of the four structures do you want - SQ-parallel, community-parallel, lifecycle-stage, or the hybrid community-body-plus-design-space-synthesis? And is the proposal's Chapter 2 the same document as the research-methodology course literature survey, or two deliverables that may legitimately have different structures?
2. Sequencing: do QL-01..QL-05 execute before or after the D-RQ-01/D-RQ-02 wording sign-off? The query concepts were frozen against the provisional wording - if 'expert judgment' replaces 'human judgment' per the open E8 choice, QL-02's concept set changes and the freeze would need re-issuing before execution rather than after.
3. Known holes in the frozen query set: the per-RQ coverage check identifies clusters that no frozen query carries terms for (learning-to-defer / selective prediction, and LLM/agent uncertainty calibration). Do we add QL-06/QL-07, or close them by documented backward/forward snowballing and accept the asymmetry in the audit trail?
4. RQ3 currently has zero tagged sources in the per-RQ map. Should Chapter 2 carry an SQ3 section that honestly states 'screened, nothing included yet', or should SQ3's related work be deferred until after the searches run?
5. Does the design-science methodology literature belong in Chapter 2 (it is currently a taxonomy branch) or in the methodology chapter? This changes the chapter boundary regardless of which structure is chosen.

**Artifact definition (§4)**

6. Artifact granularity: does each study need exactly one named artifact, or is a package acceptable? Study 3's 'Core artifact' row currently lists ten items, at least six of which are outputs of running the study rather than designed artifacts - should that row be re-cut before Chapter 4 is written?
7. Abstraction level: does the committee expect an architecture-plus-instantiation, or a higher-abstraction contribution (stated design principles, or a specification others could implement and be tested against)? This decides what section 4 is even about, and the answer differs per study.
8. SQ2/SQ3 boundary: the reuse_scope field lives in the SQ2 representation while the domain-specific vs broadly-transferable classification is described as the analytic core of SQ3. Where exactly is the line, so Studies 2 and 3 do not ship the same artifact twice? The 2026-08-10 brainstorm flagged this as unresolved and it is still unresolved.

**Evidence admissibility**

9. Admissible evidence ahead of EXP-005: is inter-rater agreement on APPLYING a designed procedure - instrument reliability, no expert gold labels - acceptable as preliminary Study 2/3 evaluation, or does every evaluation wait for at least 20 generalization-safe adjudicated labels? EXP-005 is 0 of 24 today.

10. Preliminary results in the proposal: may the offline replay series (EXP-006 event reconstruction, EXP-007 dosage-mode load-vs-coverage, EXP-008 unstable-but-never-reviewed candidates) appear in section 4 as preliminary mechanism/observability results, and with exactly what wording - in particular EXP-007's per-mode routed-item counts, which must not be phrased as an effort-reduction result?
11. Synthetic fixtures: EXP-009/EXP-010 remain provisional and gated on M-04 protocol approval. Do they appear in the proposal at all, and under what label - or are they withheld until M-04 is recorded?

**Scope, resourcing and permission**

12. Plan A presence in section 4: with medical entry gates at 0/6 and the internal 2026-08-26 checkpoint ahead, should section 4 be written Plan-B-first with the healthcare instantiation as a clearly conditional appendix, or should both paths be developed in parallel at proposal stage?
13. Study 3 resourcing: two independent reviewers plus an adjudicator are named as a hard dependency but no people are identified. Does the supervision team supply them, or does the proposal record this as an open resourcing risk with an owner and a date?
14. Permission and timing: when does the instruction to think about but not start sections 2 and 4 lift, and what is the trigger - wording sign-off, search execution, or a calendar date? Who owns the first formal section 4 draft, and against which of the artifact options above?

---

## What is deliberately NOT in this note

- No chapter outline committed for §2, and no section text drafted.
- No artifact selected for any study, and no component list frozen.
- No literature search executed; no screening, no inclusion decisions, no novelty statement.
- No evaluation design specified — that is Chapter 4 (methodology) and is out of scope for this pass.

*Prepared for the 2026-08-12 supervisor meeting. Options derive from the three-study contract, the frozen search protocol, the existing taxonomy, and the offline experiment record; the boundary facts above are unchanged by anything in this note.*